

An introduction to x86-64 assembly language programming from a compiler writer's perspective

1

What is x86-64?

- We'll be using the "x86-64" architecture as our main target:
 - A "64-bit" instruction set
 - First introduced by AMD in around 2000 as an extension of Intel's 32-bit "IA-32" or "x86" instruction set
 - Also known as "AMD64", "Intel64", "x86_64", ...
 - Broadly adopted by:
 - processors from Intel, AMD, Via, ...
 - laptops, desktops, servers, gaming consoles, ...
 - Linux, Mac OS X, Windows, ...

2

Other architectures:

- Not to be confused with:
 - IA64: a completely different 64-bit Intel architecture (Itanium)
 - ARM: widely used in phones, tablets, and more
 - IBM Power: used in Xbox 360, PS3, Wii, servers, and more
 - SPARC: used by some of the college's Unix servers
- You won't be able to run x86-64 code directly on a computer that uses one of these alternative instruction sets!

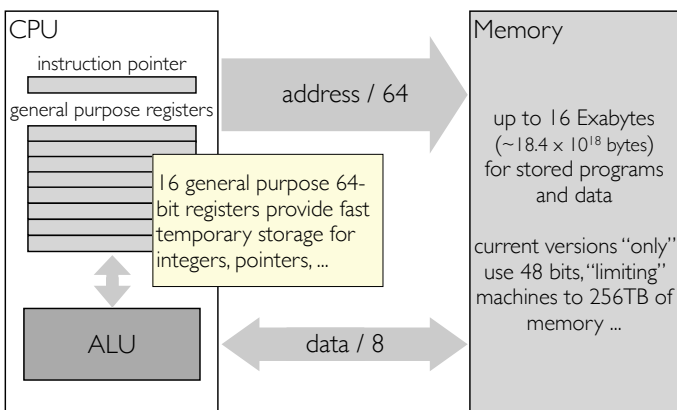
3

Notes

- No prior or in-depth knowledge of x86-64 programming will be assumed
- We will only use a small subset of the full instruction set
- If you're looking to become an expert on x86-64 programming, you'll want to look for another class!
- We'll be using the *AT&T syntax* for x86-64 assembly language rather than the *Intel syntax*. This is the default syntax used by the free GNU tools in Linux, MacOS, and DJGPP or Cygwin on Windows, and others
- For simplicity, I recommend: [ssh yourid@linuxlab.cs.pdx.edu](mailto:ssh.yourid@linuxlab.cs.pdx.edu) (or, on Windows, the equivalent using PuTTY)

4

A (greatly) simplified view of the x86-64



5

Programming for x86-64

- In concrete terms, an x86-64 program is just a collection of byte values (*machine code*)
- Once it has been loaded in to memory, the processor can *execute* a program by interpreting the byte values as *instructions* for the processor to act on
- For practical purposes, we will usually write x86-64 programs in a textual format called *assembly language* that is easier to read than the raw byte values
- A compiler that translates assembly language programs in to machine code is called an *assembler*

6

An assembly code listing

	<code>.globl f</code>	
0000 55	<code>f: pushl %ebp</code>	
0001 89E5	<code>movl %esp,%ebp</code>	
0003 53	<code>pushl %ebx</code>	
0004 8B5D08	<code>movl 8(%ebp), %ebx</code>	
0007 B8000000	<code>movl \$0, %eax # initialize length count in eax</code>	
000c EB04	<code>jmp test</code>	
000e 40	<code>loop: incl %eax # increment count</code>	
000f 83C304	<code>addl \$4, %ebx # and move to next array element</code>	
0012 8B0B	<code>test: movl (%ebx), %ecx # load array element</code>	
0014 83F900	<code>cmpl \$0, %ecx # test for end of array</code>	
0017 75F5	<code>jne loop # repeat if we're not done ...</code>	
0019 5B	<code>popl %ebx</code>	
001a 89EC	<code>movl %ebp,%esp</code>	
001c 5D	<code>popl %ebp</code>	
001d C3	<code>ret</code>	
Machine code	Assembly code	

7

An assembly code listing


addresses / offsets	label	directive	comments
0000 55	f:	<code>.globl f</code>	
0001 89E5		<code>pushl %ebp</code>	
0003 53		<code>movl %esp,%ebp</code>	
0004 8B5D08		<code>pushl %ebx</code>	
0007 B8000000		<code>movl \$0, %eax # initialize length count in eax</code>	
000c EB04		<code>jmp test</code>	
000e 40	loop:	<code>incl %eax # increment count</code>	
000f 83C304		<code>addl \$4, %ebx # and move to next array element</code>	
0012 8B0B	test:	<code>movl (%ebx), %ecx # load array element</code>	
0014 83F900		<code>cmpl \$0, %ecx # test for end of array</code>	
0017 75F5		<code>jne loop # repeat if we're not done ...</code>	
0019 5B		<code>popl %ebx</code>	
001a 89EC		<code>movl %ebp,%esp</code>	
001c 5D		<code>popl %ebp</code>	
001d C3		<code>ret</code>	
machine code	instructions		

8

x86-64 registers

9

8-bit registers (holding a single byte, 0-255)



Introduced in 1978 as part of the 8086 architecture

10

16-bit registers ("word")

accumulator	ax	ah	al
base	bx	bh	bl
count	cx	ch	cl
data	dx	dh	dl
source index	si		
destination index	di		
base pointer	bp		
stack pointer	sp		

Introduced in 1978 as part of the 8086 architecture

11

32-bit registers ("double word")

accumulator	eax	ax	ah	al
base	ebx	bx	bh	bl
count	ecx	cx	ch	cl
data	edx	dx	dh	dl
source index	esi	si		
destination index	edi	di		
base pointer	ebp	bp		
stack pointer	esp	sp		

"e" for extended
sometimes referred to as "long word"s

Introduced in 1985 as part of the 80386 architecture

12

64-bit registers (“quad word”)

rax	eax	ax	ah	al
rbx	ebx	bx	bh	bl
rcx	ecx	cx	ch	cl
rdx	edx	dx	dh	dl
rsi	esi	si		
rdi	edi	di		
rbp	ebp	bp		
rsp	esp	sp		
r8				
r9				
r10				
r11				
r12				
r13				
r14				
r15				

“r” for register

Introduced in 2000 as part of the x86-64 architecture

13

64-bit registers (“quad word”)

rax	eax	ax	ah	al
rbx	ebx	bx	bh	bl
rcx	ecx	cx	ch	cl
rdx	edx	dx	dh	dl
rsi	esi	si		sil
rdi	edi	di		dil
rbp	ebp	bp		bpl
rsp	esp	sp		spl
r8	r8d	r8w		r8b
r9	r9d	r9w		r9b
r10	r10d	r10w		r10b
r11	r11d	r11w		r11b
r12	r12d	r12w		r12b
r13	r13d	r13w		r13b
r14	r14d	r14w		r14b
r15	r15d	r15w		r15b

14

The 32-bit and 64-bit registers

rax	eax
rbx	ebx
rcx	ecx
rdx	edx
rsi	esi
rdi	edi
rbp	ebp
rsp	esp
r8	
r9	
r10	
r11	
r12	
r13	
r14	
r15	

We'll mostly just be using 64-bit and 32-bit registers

- 32-bit registers for C-style `int` values
- 64-bit registers for addresses and C-style `long` values

15

Special vs. general purpose registers

- `rip`: the instruction pointer register
- `rsp`: the stack pointer register
- `eflags`: the flags register, stores information about the results of the most recent arithmetic or logic instruction
- Other registers can typically be used for any purpose (although some instructions—division, for example—work only with specific registers)

16

x86-64 instructions

17

Instruction format

- A typical x86-64 instruction has the form:



- A suffix on the opcode indicates the size of the data that is being operated on:
 - 64-bit values use the suffix **q**(uad word)
 - 32-bit values use the suffix **l**(ong)
 - 16-bit values use the suffix **w**(ord)
 - 8-bit values use the suffix **b**(yte)

18

Addressing modes

- **Register access**, reg:
 - `%rax`: the value in register `rax`
 - Can typically use any registers except `rip` and `eflags`
- **Memory access**, mem:
 - `var`: the value in the memory location at address `var`
 - `(%rax)`: the value in memory at the address in `rax`
 - `8(%rax)`: the value in memory at the address given by adding 8 to the value in `rax`
- **Immediate**, immed:
 - `$42`: the constant value 42 (decimal; use `$0x2A` for hex)
 - `$var`: the address of memory location `var`

19

Directives for “declaring” variables

```
.data                # put variables in the “data” section
                    # (code usually goes in .text)

.align 4             # make sure address is multiple of 4
myvar: .long 42       # Simple variable, initialized to 42

.global days         # A globally accessible array of ints
days: .long 31, 28, 31, 30, 30, 30
      .long 31, 31, 30, 31, 30, 31


scratch:.space 4*100 # reserve uninitialized space

big:    .quad 123     # a 64-bit integer (takes 8 bytes)
medium: .long 123     # a 32-bit integer (takes 4 bytes)
regular:.short 123    # a 16-bit integer (takes 2 bytes)
small:  .byte 123     # an 8-bit integer (takes 1 byte)
```

20

How values are stored in memory

- A double word holds 32 binary digits (“bits”) (i.e., 4 bytes)



- `0xBE1A3910` can be interpreted as -1,105,577,712 (signed) or 3,189,389,584 (unsigned)
- Stored in memory with the least significant byte at the lowest address (“little endian”):

stored byte	0x10	0x39	0x1A	0xBE
address	400	401	402	403

21

x86-64 instructions: data movement

22

Move instructions

- Copy data from a source to a destination (where X is one of the size suffixes: b,w,l,q):

`movX src, dst`

- Any of the following combinations of arguments is allowed:

`movX reg, (reg | mem)`

`movX mem, reg`

`movX immed, (reg | mem)`

- Note that you can’t move mem to mem in one instruction

23

Examples

Suppose that the memory (starting at address 0) contains the following (eight byte) values:

8	6	2	8	0	2	4	1	7	3	4	5	6
0	8	16	24	32	40	48	56	64	72	80	88	96


Then

instruction	contents of rax
<code>movq \$24, %rax</code>	24
<code>movq (%rax), %rax</code>	8
<code>movq 24(%rax), %rax</code>	0

24


Zero and sign-extension

- Suppose we want to copy a value from a 16-bit register in to a 32-bit register



- Two common strategies:

- Zero extension: for unsigned values



- Sign extension: for signed values



25

Move with sign extension

- Copy from source to larger destination with sign extension:

`movsFT src, dst`

- F and T are the “from” and “to” sizes (either b, w, l, or q)

- Valid combinations: bw, bl, bq, wl, wq, or lq

- Examples:

```
movsbq %al, %rdi    # byte to quad
movswl %ax, %eax    # word to long
movslq %eax, %rdx    # long to quad
```

26

Move with zero extension

- Copy from source to larger destination with zero extension:

`movzFT src, dst`

- Here, FT is one of the combinations: bw, bl, bq, wl, or wq

- There is no case for lq:

- In x86-64, every instruction with a 32-bit register destination automatically zero extends the result to fill the associated 64-bit register

- For example: `movl %eax, %ebx` sets
 - the lower 32 bits of `rbx` to the value from `eax`
 - the upper 32 bits of `rbx` to zero

27

Scaled indexed addressing

- `[base] ([reg1], reg2 [, index])`

a memory operand whose address is the value in `reg1`, plus the specified base constant, plus the value of `reg2` times the index (which must be 1, 2, 4, or 8)

- Any of the parts in [...] can be omitted

- Examples:

```
(rax, rbx, 4)  the rbxth element in the array of 32-bit
               words starting at the address in rax
days(, rbx, 4) the rbxth element in the array of 32-bit
               words starting at the address days
```

28

More examples

Suppose that the memory (starting at address 0) contains the following (eight byte) values:

8	6	2	8	0	2	4	1	7	3	4	5	6
0	8	16	24	32	40	48	56	64	72	80	88	96

Then

instruction	rax	rbx
<code>movq \$24, %rax</code>	24	
<code>movq 16(%rax), %rbx</code>	24	2
<code>movq 32(%rax, %rbx, 4), %rax</code>	7	2

29

The lea (load effective address) instruction

- Load the address of the source operand (must be memory) to a destination (where X is one of the size suffixes: b, w, l, q):

`leax src, dst`

- Can also be used to co-opt the addressing mode circuitry into performing arithmetic operations:

```
leaq 4(%rax), %rax    # rax += 4
leal 1(%eax, %eax, 2), %eax # eax = 3*eax + 1
leaq 1(%rax, %rax), %rax # rax = 2*rax + 1
leal 4(, %eax, 8), %eax  # eax = 8*eax + 4
```

- These instructions just do an address calculation and do not attempt to read the data at that address.

30

The exchange instruction

- Exchange data between two locations

`xchgX (reg | mem), reg`

- Consider the following instructions in a high-level language:

```
int tmp = x;
x       = y;
y       = tmp;
```

- If `x` and `y` are held in registers, then a “clever enough” compiler can translate this code into a single `xchgl` instruction

31

The instruction pointer, `rip`

- The `rip` register holds the address of the next instruction to be executed
- As the processor reads each instruction, it increments the value in `rip` by the appropriate number of bytes to point to the following instruction
- This mechanism allows the processor to execute a sequence of instructions stored in contiguous locations in memory
- What would happen if we “move” a different value in to `rip`?

32

Jumping and labels

- We can transfer control and start executing instructions at address `addr` by using a jump instruction

`jmp addr`

- Labels can be attached to instructions in an assembly language program:

```
▶ jmp b
a:  jmp c
b:  jmp a
c:  ...
```

- Modern, pipelined machines work well with sequences of instructions that appear in consecutive locations. Jumps can be expensive: one of the goals of an optimizing compiler is to avoid unnecessary jumps.

33

x86-64 instructions: arithmetic and logic operations

34

Arithmetic instructions

- Combine a given `src` with a given `dst` value and leave the result in `dst`:

```
▶ addX  src, dst } integer arithmetic
  subX  src, dst } (signed)
  imulX src, dst }
  andX  src, dst } bitwise arithmetic
  orX   src, dst }
  xorX  src, dst }
```

- Similar to `dst += src`, `dst -= src`, etc.. in C/C++

35

Examples

- To compute $x^2 + y^2$ and store the result in `z`:

```
▶ movq x, %rax
  imulq %rax, %rax
  movq y, %rbx
  imulq %rbx, %rbx
  addq %rbx, %rax
  movq %rax, z
```

register	contents
<code>rax</code>	$x^2 + y^2$
<code>rbx</code>	y^2

```
        .data
x:      .quad 4
y:      .quad 3
```


36

x86-64 instructions: conditional execution

37

Flags

- In addition to performing the required operation, arithmetic instructions also change bits in the `eflags` register



- The flags record details about the last operation, such as:
 - Was the result zero?
 - Was the result positive?
 - Did a carry occur?
 - etc...

38

Conditional jumps, `jCC`

We can test these flags in *conditional jump* instructions

<code>jz addr</code>	(jump to <code>addr</code> if the zero flag is set)	
<code>jnz addr</code>	(jump to <code>addr</code> if the zero flag is not set)	
<code>je addr</code>	(jump to <code>addr</code> if equal; same as <code>jz</code>)	
<code>jne addr</code>	(jump to <code>addr</code> if not equal; same as <code>jnz</code>)	
<code>jle addr</code>	(jump to <code>addr</code> if less than)	} (signed)
<code>jnl addr</code>	(jump to <code>addr</code> if not less than)	
<code>jg addr</code>	(jump to <code>addr</code> if greater than)	
<code>jng addr</code>	(jump to <code>addr</code> if not greater than)	
...		

39

Examples

<code>subl %eax,%ebx</code>	jump to <code>addr</code> if <code>ebx = eax</code>
<code>jz addr</code>	
<code>subq %rax,%rbx</code>	jump to <code>addr</code> if <code>rbx ≠ rax</code>
<code>jnz addr</code>	
<code>subl %eax,%ebx</code>	jump to <code>addr</code> if <code>ebx < eax</code>
<code>jle addr</code>	
<code>subq %rax,%rbx</code>	jump to <code>addr</code> if <code>rbx ≥ rax</code>
<code>jnl addr</code>	

If the specified condition does not apply, then execution just continues with the next instruction ...

40

The compare instruction

- The `cmpX` instruction behaves like `subX` except that the result is not saved; only the flags are changed

• For example: `cmpl %eax,%ebx`
`jle addr`

will jump to `addr` if the value in `ebx` is less than the value in `eax`, but it will not change the values in either register

41

Other conditional instructions

- There are some other instructions that perform an action based on the conditional flags without the cost of a jump
- `setCC reg8` sets the value in a specified 8-bit register to 0 or 1, based on the condition specified by CC:

```
cmpl    %ecx,%ebx    # set eax to 1 if
setl     %al          # ebx < ecx, or
movzbl   %al,%eax     # else to 0
```

- `cmovCC src, dst` copies data from the specified `src` to `dst`, but only if the condition specified by CC holds:

```
cmpl     %ebx,%eax    # set eax to the max of
cmovl     %ebx,%eax    # eax and ebx
```

← condition code; no size suffix here!

42

x86-64 instructions: more arithmetic

43

Unary operations

- The following arithmetic operations have only one argument (which serves as both source and destination)

► negX	(reg mem)	negate
notX	(reg mem)	complement
incX	(reg mem)	increment
decX	(reg mem)	decrement

- Like the binary operators, these instructions also set the flags for subsequent testing


44

Division

- Divide implicit destination (`edx:eax`) (a 64-bit quantity) by a specified argument with result in `eax` and remainder in `edx`

idivl (reg | mem)

- Often used in conjunction with the `cldq` instruction (“convert long to double”, a.k.a. `cdq`), which converts a signed 32-bit value in `eax` into the corresponding signed 64-bit value in `edx:eax`.



45

Example 1

Divide 4,660 (i.e., 0x1234) by 25:

```
► movl    $0x1234, %eax
   cldq
   movl    $25, %ecx
   idivl   %ecx
```

Results: `eax` = 0xBA (186)
`edx` = 0xA (10)

Sure enough: $186 \cdot 25 + 10 = 4,660$

46

Example 2

Divide -1,105,577,712 (i.e., 0xBE1A3910) by 256

```
► movl    $0xBE1A3910, %eax
   cldq
   movl    $256, %ecx
   idivl   %ecx
```

Results: `eax` = 0xFFBE1A3A (-4,318,662)
`edx` = 0xfffff10 (-240)

Sure enough: $-4,318,662 \cdot 256 - 240 = -1,105,577,712$

47

Complications of division

- Division produces multiple results: a quotient and a remainder
- Division uses special registers: we'd better not store any other values in `eax` or `edx` if there's a chance that a division instruction might be executed
- Division can raise an exception if the src is zero (or -1)
- (We can do division on 64-bit registers in the analogous way using `rax` and `rdx`. I described the 32-bit case here because the numbers are smaller, and because that is what we'll use for `int` arithmetic.)

48

x86-64 instructions: using the stack

49


Stack

- The x86-64 includes features that allow the programmer to use a region of memory as a simple stack:
 - the `rsp` (stack pointer) register
 - special instructions like `push`, `pop`, `call`, `ret`, ...
- There is no obligation for the programmer to use these features, but it is often convenient to do so:
 - for temporary/scratch storage when a calculation needs more storage than the CPU registers can provide
 - to support calling and returning from functions

50

A typical memory layout

- A typical operating system reserves an area of scratch memory for each program, and sets the `rsp` register to point to the end of this region when the program begins



- The stack pointer moves
 - down (decreases) as values are pushed on to the stack
 - up (increases) as values are popped off of the stack
- So long as they never overlap, the data and stack areas can grow or shrink as necessary as the program runs

51

Stack operations

- Push a value onto the stack

`pushq (reg | mem | immed)`

- Pop a value of the stack

`popq (reg | mem)`

- Roughly speaking:

```
• pushq src  =  subq $8, %rsp;  movq src, (%rsp)
• popq dst   =  movq (%rsp), dst; addq $8, %rsp
```

52

Spilling temporaries on the stack

- The stack is often used for saving the contents of a register on the stack ("spilling") so that the register can be used, temporarily, for some other reason

• For example:

```
pushq %rax
pushq %rdx
... code that changes rax and/or rdx ...
popq  %rdx
popq  %rax
```

pop values in reverse order
that was used to push them!

- Note that values on the stack can still be accessed, from memory, using `(%rsp)`, `8(%rsp)`, `16(%rsp)`, `24(%rsp)`, ...

53

Call and return

- There is a special instruction for calling a function

```
call addr      =      pushq $lab
                  jmp     addr
lab: ...
```

- And a special instruction for returning from a function

```
ret            =      popq  %rax ←
                  jmp     *%rax
```

assuming
rax isn't being
used for
something
else ...

- In practice, additional instructions are often needed to deal with parameter passing, etc. ...

- (to be continued!)

special syntax: jump
to the address given
by the contents of
rax

54

Closing thoughts

55

CISC: a complex instruction set computer?

- The x86-64 is often referred to as a CISC because a single instruction can execute multiple low-level operations
- For example:

```
movq    8(%rbx, %rax, 4), %rax
does the same as:  imulq   $4, %rax
                  addq    %rbx, %rax
                  addq    $8, %rax
                  movq    (%rax), %rax
```
- Read the technical documentation to determine which sequence is faster/takes fewer bytes ... but the single instruction is usually the winner
- Using CISC instructions effectively is a major challenge for compiler writers!

56

RISC vs CISC

- RISC machines are “reduced instruction set computers” that require multiple instructions to simulate one CISC instruction ... but each RISC instruction can potentially run faster
- Easier compiler targets? RISC machines are typically simpler and more regular than CISC machines. They often have more registers that the programmer/compiler can use.
- Harder compiler targets? A general philosophy of some RISC machine designs is to let compilers rather than CPUs handle the complex tasks.

57

The GNU assembler, as

- Assembly code goes in files with a .s suffix
- We will typically use gcc to invoke the assembler

```
gcc -o output assemblyCode.s runtime.c
```
- You can also invoke the assembler directly; detailed documentation is available from:
<http://sourceware.org/binutils/docs/as/>
For x86 programming, look in particular at the section on “80386 Dependent Features”
- This information is provided only for the curious; for the purposes of this class, all the information that you need should be on these slides, or in the distributed source code.

58

Summary

- An x86-64 machine has
 - A fixed set of registers
 - Instructions for moving and operating on data
 - Instructions for testing and control transfer
- To generate good code for an x86-64 machine, we need:
 - To select appropriate instructions
 - To make good use of registers and avoid unnecessary memory accesses
 - To avoid using registers that are already in use
 - To keep the number of jumps as low as possible

59